

Security Analysis Tools for Ethereum Smart Contracts: A Comparison Based on Real-World Exploits and Vulnerabilities.

Master's Thesis by

Michele Massetti

at the KIT Department of Informatics
Institute of Information Security and Dependability (KASTEL)

Reviewer: Prof. Bernhard Beckert, Prof. Valentina Gatteschi
Advisor: Jonas Schiff

15 April 2022 – xx MONTH 20XX

Karlsruher Institut für Technologie
KIT-Fakultät für Informatik
Postfach 6980
76128 Karlsruhe

I hereby declare that the work presented in this thesis is entirely my own. I confirm that I specified all employed auxiliary resources and clearly acknowledged anything taken verbatim or with changes from other sources. I further declare that I prepared this thesis in accordance with the rules for safeguarding good scientific practice at Karlsruhe Institute of Technology (KIT).

Karlsruhe, xx MONTH 20XX

.....
(Michele Massetti)

Abstract

English abstract.

Contents

List of Figures

List of Tables

1 Introduction

1.1 Motivation

Blockchain represents one of the most popular trends in finance and computer science, during the last few years the number of investments has been growing exponentially. According CoinGeko, the crypto market's value is standing around \$2 trillion.

Bitcoin can be considered the “father” of this technology. **Bitcoin** depicted that in his paper, and in the early 2009, it was effectively launched and the cryptocurrency Bitcoin was introduced. CoinGeko states the value of Bitcoin around \$38,553.70 and its market capitalization more than \$700 billions.

Many blockchain systems have been born with new capabilities, which have allowed them to fit many different use cases. The first, which allowed developers to code on top of itself, was Ethereum. **Ethereum** published its whitepaper in 2014, and in 2015 it was deployed. The revolutionary aspect of Ethereum is the introduction of Smart Contract. These are programs running on blockchain systems and give the developers the opportunity to interact directly with this new technology. The development of innovative and prominent applications is a consequence of their development, such as NFT marketplaces, music royalty tracking, supply chain and logistics monitoring, voting mechanism, cross-border payments, and many others.

Interest in such a market has grown even among malicious attackers. Attacks such as the “Parity Wallet Hack” and the “Decentralized Autonomous Organization Attack” cost millions of dollars simply because of naive bugs in the smart contract code. Blockchain and smart contract technologies have multiple aims, but unfortunately, new applications based on them still contain bugs and multiple vulnerabilities, which cause several issues for the end-users. Most of the use of this technology relates to finance or certifications, therefore integrity, authentication and authorisation in transactions are mandatory. The research field behind blockchain technology is growing, as well as the one concerning its security and accordingly, many analysis tools were developed. These incorporate various strategies for performing the analyses, concerning the technical aspects of smart contracts, so these would work differently according to the object of the analysis.

The topic that will be addressed in this thesis work is the analysis of smart contract security properties with the usage of tools. It involves the understanding of smart contracts properties and the comparison between different tools, providing insight regarding their behaviours in different contexts.

1.2 Research Goals

Research Question: How do state-of-the-art analysis tools for Ethereum/Solidity perform (on different classes of properties/bugs)?

This thesis focuses on a dozen analytic tools, which we choose based on the type of analysis, trying to have a range of different typologies. We will test them on vulnerable smart contracts and figure out which properties are violated during real-world exploits. Furthermore, we are going to compare the tools, based on their performance, in particular, the criteria for the evaluation can cover the completeness of the analysis, the amount of found vulnerabilities and the number of false positive and negative. The execution time is crucial too, we want to understand how long it takes for finding a vulnerability. The time for the configuration and the report interpretability are parameters for defining how much a tool is user friendly.

For answering the research question, we will give an answer to sub questions such as:

- How does a tool perform the analysis?
- Which properties are relevant for smart contract security?
- Which ones have been violated in real-world exploits?
- Which tools detect which class of vulnerabilities?

1.3 Research Approach

Simil exposè

1.4 Releted Works

Simil exposè

2 Preliminary Knowledge

2.1 Blockchain

2.2 Ethereum

2.3 Smart Contract

2.4 Security Analysis

- Fuzzing
- Symbolic Execution
- Formal Verification
-

3 Real world Exploits

In this part of thesis, we introduce the objectives of the analyses: the real world exploits.

?? illustrates the attacks involve in our work, providing a brief description of those.

This thesis involves real-world exploits that have been occurred two years ago, so 2020. Our choice tries to cover different aspects and the trend of the attacks. Our work involves the security of smart contracts, written in solidity. For having a wider overview of the topic, we decided not just to provide smart contracts deployed on the Ethereum chain, but even the ones deployed on other chains, specifically Avalanche and Binance Smart Chain. Those are still EVM-compatible blockchains.

The picked attacks deal with significant projects, involving a great number of investments. The targets of attacks are not just autonomous money makers (AMM), but they can involve the implementation of a token or NFTs.

This chapter contains a subsection per each exploits. The general structure of those involves the following three main points:

1. the explanation of the protocol;
2. the exploit;
3. the properties involved in the attack.

The first part gives a general overview of the smart contract, defining its goals and main characteristics. After that, we present the exploit, illustrating the vulnerable part of the source code and the adopted strategies by the attackers. In conclusion, the depicted properties are involved in drawing the specification of the tools.

Table 3.1: Real World Exploits

Real world exploits	Detail
Aku	Funds stuck in the contract
Cover	Wrong update of data
BZX	Wrong implementation of token
Spartan	Bad estimation of rewards in a pool
Uranium	Vulnerable deposit withdraw process in a pool
XSURGE	Reentrancy in a token
BurgerSwap	Reentrancy in an AMM
DirtyDogs	Reentrancy in NFTs market

3.1 \$34 Million stacks NFT Project Aku Dreams Smart Contract

Business2community estimates the value of NFT market around \$100 billion. Nowadays, the word NFT is one of the most researched ones on Google and the other search engine. NFT's marketplaces manage the transaction behind these valuable markets. They are made by a frontend part, but even by a backend one which relies on the blockchain. Akutars, a highly anticipated Ethereum-based NFT project developed by Aku Dreams, is an example of how a bug can have catastrophic consequences in this sector.

3.1.1 Akutarts NFT project

[Aku] reports Akutarts locked up \$34 million due to the faulty code of the smart contract. The launch contained 15,000 NFTs and was based on the Dutch auction. This strategy involves a descending price auction where an item begins at a set maximum price. The price is gradually lowered over a fixed time until a bid is placed that guarantees the bidder the purchase of the item at the current price. Anyone who paid the higher amount would get a refund. Unfortunately, the launch was corrupted, since the errors in the codes made the project open to exploits. An attacker could block the withdrawals and refunds while attempting to highlight the vulnerabilities within the project.

3.1.2 The exploit

The first part of the exploit involved the function processRefunds ??.

This has the aim to refund the bid of the user who took part in the auction.

The problem relies on the for loop in line 11. It loops on all over the users, who needs to be refunded, estimating the number of tokens to send. Then, the amount is sent with the function call, which returns a boolean, based on the correct execution of the operation. This is checked with a require, line 22; so, if the operation concluded incorrectly, it would revert the transaction.

The problem relies on the require in the loop. If one of the accounts could not receive the refund, the function would always reverted. Since looping all over the users is a sequential operation, if the transaction just reverted when it reaches an item, it would never reach all the following items.

Therefore, a malicious user just implemented a smart contract which took part in the auction and reverted any time it received tokens.

```
1 function processRefunds() external {
2     require(block.timestamp > expiresAt, "Auction still in progress");
3     uint256 _refundProgress = refundProgress;
4     uint256 _bidIndex = bidIndex;
5     require(_refundProgress < _bidIndex, "Refunds already processed");
6
7     uint256 gasUsed;
8     uint256 gasLeft = gasleft();
9     uint256 price = getPrice();
10
11     for (uint256 i=_refundProgress; gasUsed < 5000000 && i < _bidIndex; i++) {
12         bids memory bidData = allBids[i];
```

```

13     if (bidData.finalProcess == 0) {
14         uint256 refund = (bidData.price - price) * bidData.bidsPlaced;
15         uint256 passes = mintPassOwner[bidData.bidder];
16         if (passes > 0) {
17             refund += mintPassDiscount * (bidData.bidsPlaced < passes ? bidData.
bidsPlaced : passes);
18         }
19         allBids[i].finalProcess = 1;
20         if (refund > 0) {
21             (bool sent, ) = bidData.bidder.call{value: refund}("");
22             require(sent, "Failed to refund bidder");
23         }
24     }
25
26     gasUsed += gasLeft - gasleft();
27     gasLeft = gasleft();
28     _refundProgress++;
29 }
30
31 refundProgress = _refundProgress;
32 }

```

Listing 3.1: Function for refunding the users.

The second part of the exploit is characterized by a bug in the logic, which could not allow the developer team to withdraw the project funds.

The function `claimProjectFunds` (??), callable only by the owner of the contract due to the modifier `onlyOwner`, refunds the developers just when all the users are considered refunded.

The boolean condition, contained in the `require` at line 3, is the heart of the problem. The `require` compares the variable `refundProgress`, which takes track of the refund progress, and `totalBids`.

```

1 function claimProjectFunds() external onlyOwner {
2     require(block.timestamp > expiresAt, "Auction still in progress");
3     require(refundProgress >= totalBids, "Refunds not yet processed");
4     require(akuNFTs.airdropProgress() >= totalBids, "Airdrop not complete");
5
6     (bool sent, ) = project.call{value: address(this).balance}("");
7     require(sent, "Failed to withdraw");
8 }

```

Listing 3.2: Function for claiming the funds for the developers.

The variable `totalBids` is increased every time a bid is placed, regardless of the user who computed it, shown in ?? at line 18. The user can call the function `_bid`, for placing a bid, with an arbitrary amount of bids, but the variable `refundProgress` is increased every time a user is refunded. Consequently, if a user bought more than one bid, the amount of refunded users would never be greater or equal to the number of placed bids.

```

1 function _bid(uint8 amount, uint256 value) internal {
2     require(block.timestamp > startAt, "Auction not started yet");
3     require(block.timestamp < expiresAt, "Auction expired");
4     uint80 price = getPrice();

```

```
5      uint256 totalPrice = price * amount;
6      if (value < totalPrice) {
7          revert("Bid not high enough");
8      }
9
10     uint256 myBidIndex = personalBids[msg.sender];
11     bids memory myBids;
12     uint256 refund;
13
14     if (myBidIndex > 0) {
15         myBids = allBids[myBidIndex];
16         refund = myBids.bidsPlaced * (myBids.price - price);
17     }
18     uint256 _totalBids = totalBids + amount;
19     myBids.bidsPlaced += amount;
20
21     if (myBids.bidsPlaced > maxBids) {
22         revert("Bidding limits exceeded");
23     }
24
25     if(_totalBids > totalForAuction) {
26         revert("Auction Full");
27     } else if (_totalBids == totalForAuction) {
28         expiresAt = block.timestamp; //Auction filled
29     }
30
31     myBids.price = price;
32
33     if (myBidIndex > 0) {
34         allBids[myBidIndex] = myBids;
35     } else {
36         myBids.bidder = msg.sender;
37         personalBids[msg.sender] = bidIndex;
38         allBids[bidIndex] = myBids;
39         bidIndex++;
40     }
41
42     totalBids = _totalBids;
43     totalBidValue += totalPrice;
44
45     refund += value - totalPrice;
46     if (refund > 0) {
47         (bool sent, ) = msg.sender.call{value: refund}("");
48         require(sent, "Failed to refund bidder");
49     }
50 }
```

Listing 3.3: Function for users'bid

3.1.3 Properties

The smart contract involves 2 main problems: the refunding of the users who placed the bids and the claim of the developers' rewards.

The first property deals with the function `processRefunds`. It reverts every time because a malicious wallet, which can't receive any tokens triggering the require. We verify if the contract can always refund all the users. The property involves the sum of refunded wallets may be equal to the number of the counter which loops on the map containing the data. It is a postcondition, so it is proven that all the users are refunded if the function does not revert, which means that the function does not consider the case of error.

The other property regards the function `claimProjectFunds`. In the beginning, some requirements have to be fulfilled before the owner can obtain the rewards. Our focus is on the comparison between the counter of the refunded users and the total amount of bids. In this case, we use proof by contradiction. We check if the `processRefunds` variable is always less than the `totalBids`. The property should be proofed if we consider that at least one user placed more than one bid.

3.2 Cover Protocol:Infinite Minting Exploit Nets Attacker \$4.4M

On the 28th of December 2020, an exploit was abused on Cover Protocol's shield mining contract. The article shows the attackers could steal from project around \$ 4 million. The target of the attack was the smart contract `Blacksmith.sol`, its bug had the result to mint more rewards to the miner.

3.2.1 Cover Protocol

CoverProtocol interviewed the co-founder of the Cover Protocol. In his article he answers some question about his project, regarding its functionality and road map. It was an active protocol on the Ethereum blockchain; the developer deployed version 2, because of the attack. Cover Protocol is a peer-to-peer coverage marketplace that utilizes ERC-20 fungible tokens to allow permissionless and non-KYC coverage. It can be described as a coverage provider. The attack affected the rewards contract, consequently, the token's one even. The exploit can be classified under the name of "infinite mint".

3.2.2 The exploit

The developports's team reported [**CoverProtocolPostMortem**] the technical analysis of the exploit the day after. The contract containing the vulnerability is `Blacksmith.sol`. The core protocol was not affected, but the minting contract and the `$COVER` token became unusable. Firstly, the attackers created a new balancer liquidity pool for the target contract. The next step was to deposit token in it and execute the exploit, withdrawing funds from the contract thanks to a miscalculation of the rewards. The bug relies on the misuse of two keywords in solidity: storage and memory.

Memory This keyword within Solidity allocates memory for a specific variable. In this instance, that variable is scoped to a specific function. The memory is cleared once the function has executed.

Storage On the other hand this keyword within Solidity allows variables to act as a pointer into the storage of data in mappings or data structures. Storage data is persistent between function calls and transactions.

The previous has a similar behave to the Random Access Memory (RAM) on a computing device, the latter stores into the persistent memory.

The vulnerable function is the deposit one.

```
1  function deposit(address _lpToken, uint256 _amount) external override {
2      require(block.timestamp >= START_TIME , "Blacksmith: not started");
3      require(_amount > 0, "Blacksmith: amount is 0");
4      Pool memory pool = pools[_lpToken];
5      require(pool.lastUpdatedAt > 0, "Blacksmith: pool does not exists");
6      require(IERC20(_lpToken).balanceOf(msg.sender) >= _amount, "Blacksmith: insufficient
    balance");
7      updatePool(_lpToken);
8
9      Miner storage miner = miners[_lpToken][msg.sender];
10     BonusToken memory bonusToken = bonusTokens[_lpToken];
11     _claimCoverRewards(pool, miner);
12     _claimBonus(bonusToken, miner);
13
14     miner.amount = miner.amount.add(_amount);
15     // update writeoff to match current acc rewards/bonus per token
16     miner.rewardWriteoff = miner.amount.mul(pool.accRewardsPerToken).div(CAL_MULTIPLIER)
    ;
17     miner.bonusWriteoff = miner.amount.mul(bonusToken.accBonusPerToken).div(
    CAL_MULTIPLIER);
18
19     IERC20(_lpToken).safeTransferFrom(msg.sender, address(this), _amount);
20     emit Deposit(msg.sender, _lpToken, _amount);
21 }
```

Listing 3.4: Deposit function.

At line 4 of ??, the state of the pool is stored in a variable with the keyword memory. The function update updates the state of the pool. However, the variable pool, existing within the function, remains identical.

Then, deposit function at line 16 ?? estimates the reward per token updating the value of miner.rewardWriteoff, but it uses the wrong value of the parameter of pool.accRewardsPerToken.

Following the vulnerability, anyone can obtain an insane amount of minted tokens when they execute the claimRewards(address _lpToken) function. This function, which is used to grab their rewards, ends up calling _claimCoverRewards(Pool memory pool, Miner memory miner) which references the miner.rewardWriteoff. As that variable is much smaller than the actual pool.accRewardsPerToken, the contract results in minting an abundance of tokens.

3.2.3 Properties

The heart of the problem is the wrong management of the keywords storage and memory.

The consequence of this error is a miscalculation of the reward of the miner. The property relies on how it is computed. It is not estimated considering the correct parameters of the pool.

The post-condition involves the estimation of the reward inside the function deposit.

We compare the miner.rewardWriteoff and then we recumpute its mathematical operation with the updated parameters of the pool:

```
miner.amount.mul(pool.accRewardsPerToken).div(CAL_MULTIPLIER).
```

The discriminating in the operation should be the parameter pool.accRewardsPerToken.

3.3 DeFi platform bZX: \$8M hack from one misplaced line of code

bZxProtocol explains how this protocol works. Anyone can use bZx to create apps that allow lenders, borrowers, and traders to interact with Ethereum based decentralised finance protocol. It is a community-run project, moreover all major protocol changes requiring a community vote.

Protocols can be developed by bZx protocol, an example is Fulcrum. It is a powerful DeFi platform for tokenized lending and margin trading. iTokens (margin loans) represent the earn holders interest on borrowed funds and pTokens (tokenized margin positions) allow your margin positions to be composable.

Unfortunately, it suffered a couple of attacks in February 2020. The developers explained the attackers could drain different currencies, 219,199.66 LINK, 4,502.70 Ether (ETH), 1,756,351.27 Tether (USDT), 1,412,048.48 USD Coin (USDC) and 667,988.62 Dai (DAI): a total of \$8 million in value. The attack depends on a bug based on an incorrect sequence of operations.

The object of the attack was the contract named LoanTokenLogicStandard. It implements the logic behind the protocol, for managing the borrows, loans and all the functionalities. Every ERC20 token has a transferFrom() function, which has the aim to transfer the tokens. Calling this function allowed the attacker to create and transfer an iToken to himself: his balance could be artificially increased. The duplicated tokens were then redeemed for their underlying collateral, with the hackers now “owning” a much higher percentage of the pool, so the attacker could withdraw the tokens.

The snipped code ?? shows the vulnerable function. The attacker called the function with the same amount of _from and _to. Since both addresses refer to the same one, line 27 decreases the balance of the address, but then line 31 increases the same balance. The problem relies on the estimating of the amount: it is the sum of the sent token and a variable (line 23), which stored the value of the balance before the transaction.

```
1 contract LoanTokenLogicStandard is AdvancedToken, GasTokenUser {
2     using SafeMath for uint256;
3     using SignedSafeMath for int256;
4
5     modifier settlesInterest() {
6         _settleInterest();
7         -;
```

```
8 }
9 ...
10 function _internalTransferFrom(
11     address _from,
12     address _to,
13     uint256 _value,
14     uint256 _allowanceAmount)
15     internal
16     returns (bool)
17 {
18     if (_allowanceAmount != uint256(-1)) {
19         allowed[_from][msg.sender] = _allowanceAmount.sub(_value, "14");
20     }
21     //Vulnerable lines
22     uint256 _balancesFrom = balances[_from];
23     uint256 _balancesTo = balances[_to];
24
25     require(_to != address(0), "15");
26
27     uint256 _balancesFromNew = _balancesFrom
28         .sub(_value, "16");
29     balances[_from] = _balancesFromNew;
30
31     uint256 _balancesToNew = _balancesTo
32         .add(_value);
33     balances[_to] = _balancesToNew;
34
35     // handle checkpoint update
36     uint256 _currentPrice = tokenPrice();
37
38     _updateCheckpoints(
39         _from,
40         _balancesFrom,
41         _balancesFromNew,
42         _currentPrice
43     );
44     _updateCheckpoints(
45         _to,
46         _balancesTo,
47         _balancesToNew,
48         _currentPrice
49     );
50
51     emit Transfer(_from, _to, _value);
52     return true;
53 }
54 ...
```

Listing 3.5: Vulnerable function in LoanTokenLogicStandard contract.

The developers corrected the bug in few days. It was enough switching some line of code, in order to avoid the operations of sum and subtraction operate on the same balance. The code ?? presents some differences. The operations regarding the receiver's balance are computed (lines 13-15), then those which deal with the sender's one (16-20).


```

1  function _internalTransferFrom(
2      address _from,
3      address _to,
4      uint256 _value,
5      uint256 _allowanceAmount)
6      internal
7      returns (bool)
8  {
9      if (_allowanceAmount != uint256(-1)) {
10         allowed[_from][msg.sender] = _allowanceAmount.sub(_value, "14");
11     }
12     require(_to != address(0), "15");
13     uint256 _balancesFrom = balances[_from];
14     uint256 _balancesFromNew = _balancesFrom
15         .sub(_value, "16");
16     balances[_from] = _balancesFromNew;
17     uint256 _balancesTo = balances[_to];
18     uint256 _balancesToNew = _balancesTo
19         .add(_value);
20     balances[_to] = _balancesToNew;
21     // handle checkpoint update
22     uint256 _currentPrice = tokenPrice();
23     _updateCheckpoints(
24         _from,
25         _balancesFrom,
26         _balancesFromNew,
27         _currentPrice
28     );
29     _updateCheckpoints(
30         _to,
31         _balancesTo,
32         _balancesToNew,
33         _currentPrice
34     );
35     emit Transfer(_from, _to, _value);
36     return true;
37 }

```

Listing 3.6: Corrected bug in LoanTokenLogicStandard contract.

3.3.1 Properties

The function `internalTransfer` is the one which contains the bug abused by the attackers. We define 2 properties for defining the correct execution of the function. Both of those are broken by the wrong implementation of the function.

The first one is a post-condition. It involves the correct estimation of the balance of the addresses involved in the operation, the parameters `from` and `to`. The balance of the sender should decrease and the one of the receiver should increase.

The other one is an invariant. It states the total sum of balances should be less than the variable total supply.

3.4 XSURGE on BSC Chain

The **XSurgeWeb**'s whitepaper provides a presentation of the ecosystem. It is described as a great DeFi investing idea based on proprietary pricing algorithms embedded in the Surge Token Variants' contracts. Surge Token Variants each have their own Market Maker, allowing them to trade continuously and outlast both centralised and decentralised exchanges. The strategy is to reward long-term holding by increasing a holder's claim of the backing asset. Each Surge Token utilizes a built-in contract exchange system that renounces the need for a traditional liquidity pool. Both assets are stored within the contract itself, rather than a liquidity pool pair of the backing asset to the token using a traditional market maker method for exchange and price calculation.

One of the Surge Token is SurgeBNB, the one which is my focus of analysis. **XSurgeBNB** explains in deep how the attack to this contract occurred. The Official claimed that the attacker had stolen \$5 million in SurgeBNB through a backdoor vulnerability. XSURGE stated that a potential security vulnerability in the SurgeBNB contract was discovered on August 16th.

The attack is made by 4 main steps:

1. the attacker borrow 10,000BNB through flash loans.
2. Use all the BNB to buy SURGE. According to the current price, the attacker can buy 1,896,594,328,449,690 SURGE
3. He calls the "sell" function, for selling the obtained SURGE.
4. The sale function alters the data after the transfer, and the transfer code has a reentrance vulnerability. When the attack contract acquires BNB, the period before the SURGE contract's state changes (References1st:SellSURGE line 15), the attack contract can use the reentrance vulnerability to purchase SURGE again.

```
1 function sell(uint256 tokenAmount) public nonReentrant returns (bool) {
2
3     address seller = msg.sender;
4
5     // make sure seller has this balance
6     require(_balances[seller] >= tokenAmount, 'cannot sell above token amount');
7
8     // calculate the sell fee from this transaction
9     uint256 tokensToSwap = tokenAmount.mul(sellFee).div(10**2);
10
11    // how much BNB are these tokens worth?
12    uint256 amountBNB = tokensToSwap.mul(calculatePrice());
13
14    // send BNB to Seller
15    (bool successful,) = payable(seller).call{value: amountBNB, gas: 40000}("");
16    if (successful) {
17        // subtract full amount from sender
18        _balances[seller] = _balances[seller].sub(tokenAmount, 'sender does not have
19        this amount to sell');
20        // if successful, remove tokens from supply
```

```

20     _totalSupply = _totalSupply.sub(tokenAmount);
21 } else {
22     revert();
23 }
24 emit Transfer(seller, address(this), tokenAmount);
25 return true;
26 }

```

Listing 3.7: Sell function of Surge (SURGE) token.

The bnb Amount of the contract stays intact, and the total amount of SURGE tokens `totalSupply` has not been updated, because the attack contract spends all of the BNB balance to acquire SURGE each time (still remains the quantity before the sell). As a result, the price of token falls, allowing the attacker to purchase additional SURGE.

```

1  function purchase(address buyer, uint256 bnbAmount) internal returns (bool) {
2      // make sure we don't buy more than the bnb in this contract
3      require(bnbAmount <= address(this).balance, 'purchase not included in balance');
4      // previous amount of BNB before we received any
5      uint256 prevBNBAmount = (address(this).balance).sub(bnbAmount);
6      // if this is the first purchase, use current balance
7      prevBNBAmount = prevBNBAmount == 0 ? address(this).balance : prevBNBAmount;
8      // find the number of tokens we should mint to keep up with the current price
9      uint256 nShouldPurchase = hyperInflatePrice ? _totalSupply.mul(bnbAmount).div(
address(this).balance) : _totalSupply.mul(bnbAmount).div(prevBNBAmount);
10     // apply our spread to tokens to inflate price relative to total supply
11     uint256 tokensToSend = nShouldPurchase.mul(spreadDivisor).div(10**2);
12     // revert if under 1
13     if (tokensToSend < 1) {
14         revert('Must Buy More Than One Surge');
15     }
16
17     // mint the tokens we need to the buyer
18     mint(buyer, tokensToSend);
19     emit Transfer(address(this), buyer, tokensToSend);
20     return true;
21 }

```

Listing 3.8: Purchase function of Surge (SURGE) token.

Repeating three times of Round 2 and Round 3, the attacker accumulates a large amount of SURGE through reentry, and then sells all the SURGE to make a profit.

At the end of this transaction, the attack contract sold 1,864,120,345,279,610,000 SURGE, obtained 10327 BNB, and finally the profitable 297 BNB was sent to the attacker's address.

The following are the modifications suggested by the Beosin technical team for this attack:

- any transfer operation should be place after the state changes to avoid reentry assaults.
- Instead of using "call. value," use transfer or send to transfer.

3.4.1 Properties

This exploit represents a typical case of reentrancy.

The attacker's strategy involves the function `sell`, which contains the bug, and then the function `purchase`. After calling the first one and triggering the reentrancy, the malicious fallback implemented by the attacker uses the amount of money for buying more XSURGE tokens. At the end of the selling process, the total supply should decrease the amount sold by the user. But since the attacker called the `purchase`, the variable is not updated as it was supposed to be. Buying the same amount of sold tokens, the value would not change.

We define the property as a postcondition, which states the variable `_totalSupply` is decreased of the amount sold by the user, then `tokenAmount`.

3.5 CBDAO: an example of rug pull

Developers should watch out for possible attacks. They should audit and test their contract to find possible vulnerabilities and apply patches. In the decentralized finance context, even the investors should worry about malicious developers, who convince the investors to invest and then steal their investments. These class of fraud are basically type of exit scam and decentralized finance (DeFi) exploit, it is classified with the name of rug pull.

RugPullDef defines rug pull as a type of crypto scam that occurs when a team pumps their project's token before disappearing with the funds, leaving their investors with a valueless asset. Fraudulent developers create a new crypto token, pump up the price and then pull as much value out of them as possible before abandoning them as their price drops to zero.

An example of this type of fraud is the one presented in the article **CBDAO**. It seems the malicious developers could steal around 1 million dollar in ethereum (ETH).

The project main token was \$BREE. For attracting early investors, they associated to it a presale token, named \$SBREE. The ones who bought that, could swap their amount of presale token in \$BREE once the token was published, having an advantage. Unfortunately, one of the admin wallets exploited a backdoor in the SBREE token contract, minted 50,000 SBREE. After that, the attacker sold that amount in BREE token and sold it on the market. That pushed down the price of BREE at the expense of other holders. The 50,000 BREE was sold for under 200 ETH.

Following the operation of the malicious developer, it is possible to understand how the fraud occurred. This transaction, achieved by etherscan, shows the attacker called the `mint` function and could generate 50,000 SBREE. After that, it called the `BreePurchase` contract for swapping the token in BREE and then swap those in ETH on Uniswap.

The backdoor relies on the malicious management of access control. The admin, with the function `grantRole`, allow another wallet to be the Minter, so it called the function `mint`.

```
1 ...
2
3 function _grantRole(bytes32 role, address account) private {
4     if (_roles[role].members.add(account)) {
5         emit RoleGranted(role, account, _msgSender());
```

```

6     }
7 }
8 ...
9
10
11 contract Roles is AccessControl {
12
13     bytes32 public constant MINTER_ROLE = keccak256("MINTER");
14     bytes32 public constant OPERATOR_ROLE = keccak256("OPERATOR");
15
16     constructor () public {
17         _setupRole(DEFAULT_ADMIN_ROLE, _msgSender());
18         _setupRole(MINTER_ROLE, _msgSender());
19         _setupRole(OPERATOR_ROLE, _msgSender());
20     }
21
22     modifier onlyMinter() {
23         require(hasRole(MINTER_ROLE, _msgSender()), "Roles: caller does not have the MINTER
24         role");
25         _;
26     }
27
28     modifier onlyOperator() {
29         require(hasRole(OPERATOR_ROLE, _msgSender()), "Roles: caller does not have the
30         OPERATOR role");
31         _;
32     }
33 }
34
35 //the contract inherit Roles contettract
36 ...
37 modifier onlyMinter() {
38     require(hasRole(MINTER_ROLE, _msgSender()), "Roles: caller does not have the MINTER
39     role");
40     _;
41 }
42 ...
43 function _mint(address account, uint256 amount) internal virtual {
44     require(account != address(0), "ERC20: mint to the zero address");
45
46     _beforeTokenTransfer(address(0), account, amount);
47
48     _totalSupply = _totalSupply.add(amount);
49     _balances[account] = _balances[account].add(amount);
50     emit Transfer(address(0), account, amount);
51 }
52 ...

```

Listing 3.9: Backdoor inside the contract

3.6 A flash loan used for amplify a bug: \$30M drained from Spartan protocol

Spartan Protocol is a DeFi protocol for synthetic assets running on BinanceSmartChain. It inherits many capabilities of UniswapV2 protocol, adapting the code for new use cases and implementing different strategies. The fee mechanism is modified to incentivize liquidity providers when liquidity is scarce. Consequently, users trading larger volumes are charged more fees. Similar to UniswapV2, pairs WBNB and SPARTA token are open for users to add/remove liquidity. For clarifying this, let's consider the following example. Bob is able to send (WBNB+SPARTA) into the WBNB-SPARTA pool and get Liquidity Pool (LP) tokens back, redeemable for the underlying assets.

This protocol was the target of an exploit at the end of May 2021. The presence of a bug inside the code, plus the amplification due to a flash loan, allowed the attacker to drain the liquidity.

The articles **FlashCoin** and **FlashCloud** give a definition flash loan.

Flash loan A flash loan is a relatively new type of uncollateralized lending that has become popular across a number of decentralized finance (DeFi) protocols based on the Ethereum network. When it has been issued, the smart contract certifies that the borrower pays back the loan before the transaction ends. If this condition is not fulfilled, the transaction reverts, consequently the amount of loan is given back.

3.6.1 The exploit

The exploit involved 2 contracts of protocol: `Utils.sol` and `poolFactory.sol`. The latter implements the strategy for the management of the liquidity in the pool and the former provides support functions. The mistake of the developers was not to consider the updated value of underlying assets. Those are stored into the variables (`baseAmount`, `tokenAmount`) and estimated with `iBEP20(token).balanceOf(pool)` and `iBEP20(base).totalSupply()`.

The bug in code lies in the `calcLiquidityShare()` function, called in `RemoveLiquidity()`.

```
1
2 function calcLiquidityShare(uint units, address token, address pool, address member)
3   public view returns (uint share){
4     // share = amount * part/total
5     // address pool = getPool(token);
6     uint amount = iBEP20(token).balanceOf(pool);
7     uint totalSupply = iBEP20(pool).totalSupply();
8     return(amount.mul(units)).div(totalSupply);
9   }
```

Listing 3.10: `calcLiquidityShare` function

It should get the balance of the underlying asset in the pool, in line 5, ???. The amount of that which should be transferred out is calculated based on the total LP tokens supplied (line 6) and the number of LP tokens to burn (units). The function does not consider who transfers assets into the pool. The value of underlying assets can be manipulated and increased by an exploit. The real values, estimated in line 5, are different from the

ones contained in the variable (baseAmount, tokenAmount). The removeLiquidity() calls calcLiquidityShare on TOKEN and BASE, line 4 and 5 (??). It fails to synchronize the balances of the underlying assets and the variables which store the amount of the assets.

```

1 // Remove Liquidity for a member
2 function removeLiquidityForMember(address member) public returns (uint outputBase, uint
  outputToken) {
3     uint units = balanceOf(member);
4     outputBase = iUTILS(_DAO().UTILS()).calcLiquidityShare(units, BASE, address(this),
  member);
5     outputToken = iUTILS(_DAO().UTILS()).calcLiquidityShare(units, TOKEN, address(this),
  member);
6     _decrementPoolBalances(outputBase, outputToken);
7     _burn(address(this), units);
8     iBEP20(BASE).transfer(member, outputBase);
9     iBEP20(TOKEN).transfer(member, outputToken);
10    emit RemoveLiquidity(member, outputBase, outputToken, units);
11    return (outputBase, outputToken);
12 }

```

Listing 3.11: Function for Removing Liquidity

As a consequence, the _decrementPoolBalance(), line 6, updates the wrong value of the variables storing the assets. It does not get the update-to-date balances of BASE and TOKEN. Instead, it only decrements the reserved amounts (baseAmount, tokenAmount). The attacker followed these steps for draining the liquidity:

1. Add liquidity and get LP tokens back.
2. Transfer some assets into the Pool contract to amplify the number of underlying assets of the LP tokens collected in step 1.
3. Remove liquidity and get more assets than what you added in Step 1.
4. Add the assets transferred into the Pool contract as liquidity and remove them immediately.

```

1 function _decrementPoolBalances(uint _baseAmount, uint _tokenAmount) internal {
2     uint _removedBase = iUTILS(_DAO().UTILS()).calcShare(_baseAmount, baseAmount,
  baseAmountPooled);
3     uint _removedToken = iUTILS(_DAO().UTILS()).calcShare(_tokenAmount, tokenAmount,
  tokenAmountPooled);
4     baseAmountPooled = baseAmountPooled.sub(_removedBase);
5     tokenAmountPooled = tokenAmountPooled.sub(_removedToken);
6     baseAmount = baseAmount.sub(_baseAmount);
7     tokenAmount = tokenAmount.sub(_tokenAmount);
8 }

```

Listing 3.12: Function which updates decrements the assets in the pool.

A solution for this bug is shown in ???. It updates the variables of assets at line 3, before it is estimating the the amount to drain.

```
1  function calcLiquidityShareSynch(uint units, address token, address pool, address member
   ) public view returns (uint share){
2      // synchronize the variable
3      iPOOL(pool).sync();
4      uint amount = iBEP20(token).balanceOf(pool);
5      uint totalSupply = iBEP20(pool).totalSupply();
6      return(amount.mul(units)).div(totalSupply);
7  }
8
9  function sync() public {
10     baseAmount = iBEP20(BASE).balanceOf(address(this));
11     tokenAmount = iBEP20(TOKEN).balanceOf(address(this));
12 }
```

Listing 3.13: Possible correct calcLiquidityShare.

3.6.2 Properties

The issue correlated to the logic of the program involves the removing liquidity process. The attacker could amplify the number of tokens to remove.

The function calcLiquidityShare is an internal function which estimates the total of underlying assets to send to the user. It miscalculates the value because it considers the balance of the account, which can be simply modified by transferring tokens into the pool. A proposed solution implies a function, synch, which synchronises the variables baseAmount and tokenAmount, which keep track of the underlying assets in the contract. Function synch should be called in the buggy function.

This one is encompassed in the definition of the property. It is called at the end of the process of the removing the liquidity for a member. The postcondition specifies that the function synch should not update the two variables.

3.7 Uranium Finance: \$1.3M of rewards drawn

Uranium Finance is a Automated Market Maker (AMM) running on the BinanceSmartChain. The article presented by **UraniumPM**, deals with the exploit which occurred on the 8th April 2021. The attacker could grab the contents of the RADS pool and all of the RADS/sRADS rewards and sell them for \$1.3M worth of BUSD and BNB.

The team of developer could identify the exploiter, because some transaction of the attacker wallet, could be correlated with a Binance wallet. The criminal got in touch with the developers. After some negotiation, the exploiter refund the team of \$1M in ETH.

3.7.1 The exploit

The article written by **UraniumTech**, gets more in depth into the technical details involved in this exploit. The target of it was the contract MasterUranium, specifically the part regarding the rewarding of the user. The list of transactions involving the malicious wallet shows the attacker could draw a huge amount of rewards by calling 3 functions multiple times:

1. deposit(_pid, _amount);
2. emergencyWithdraw(_pid);
3. withdraw(_pid, _amount).

Deposit The two most relevant variables to the exploit are user.amountWithBonus and user.rewardDebt, for the attack purpose, they need to be greater than 0. Therefore this function is called with the _amount input argument larger than “0”. the “_bonusAmount” is calculated with:

`_bonusAmount=_amount.mul(userBonus(_pid, _user).add(10000)).div(10000).`

The user.amountWithBonus increases by adding the _bonusAmount. The user.rewardDebt is calculated by the end of the function, with `user.rewardDebt = user.amountWithBonus.mul(pool.accRadsP` When the function returns, the both variables are greater than 0.

```

1  function deposit(uint256 _pid, uint256 _amount) external validatePool(_pid) {
2      address _user = msg.sender;
3      PoolInfo storage pool = poolInfo[_pid];
4      UserInfo storage user = userInfo[_pid][_user];
5      updatePool(_pid);
6      if (user.amount > 0) {
7          uint256 pending = user.amountWithBonus.mul(pool.accRadsPerShare).div(1e12).sub(
            user.rewardDebt);
8          if(pending > 0) {
9              if(pool.isSRadsRewards){
10                 safeSRadsTransfer(_user, pending);
11             }
12             else{
13                 safeRadsTransfer(_user, pending);
14             }
15         }
16     }
17     if (_amount > 0) {
18         pool.lpToken.safeTransferFrom(address(_user), address(this), _amount);
19         if (address(pool.lpToken) == address(rads)) {
20             uint256 transferTax = _amount.mul(2).div(100);
21             _amount = _amount.sub(transferTax);
22         }
23         if (pool.depositFeeBP > 0) {
24             uint256 depositFee = _amount.mul(pool.depositFeeBP).div(10000);
25             pool.lpToken.safeTransfer(feeAddress, depositFee);
26             user.amount = user.amount.add(_amount).sub(depositFee);
27             uint256 _bonusAmount = _amount.sub(depositFee).mul(userBonus(_pid, _user).
            add(10000)).div(10000);
28             user.amountWithBonus = user.amountWithBonus.add(_bonusAmount);
29             pool.lpSupply = pool.lpSupply.add(_bonusAmount);
30         } else {
31             user.amount = user.amount.add(_amount);
32             uint256 _bonusAmount = _amount.mul(userBonus(_pid, _user).add(10000)).div
            (10000);
33             user.amountWithBonus = user.amountWithBonus.add(_bonusAmount);
34             pool.lpSupply = pool.lpSupply.add(_bonusAmount);

```

```
35     }
36   }
37   user.rewardDebt = user.amountWithBonus.mul(pool.accRadsPerShare).div(1e12);
38   emit Deposit(_user, _pid, _amount);
39 }
40
41 // Withdraw LP tokens from MasterUranium.
```

Listing 3.14: Deposit Function

EmergencyWithdraw The next step is the withdrawal of the funds. This function has the purpose of getting the deposited token back and setting user.amount equal to and user.rewardDebt equal to 0. The fundamental variable user.amountWithBonus is still larger than 0. It is exploited during the last step.

```
1 // Withdraw without caring about rewards. EMERGENCY ONLY.
2 function emergencyWithdraw(uint256 _pid) external {
3   PoolInfo storage pool = poolInfo[_pid];
4   UserInfo storage user = userInfo[_pid][msg.sender];
5   pool.lpToken.safeTransfer(address(msg.sender), user.amount);
6   emit EmergencyWithdraw(msg.sender, _pid, user.amount);
7   user.amount = 0;
8   user.rewardDebt = 0;
9 }
```

Listing 3.15: Deposit Function

Withdraw In the last step, the attacker call this function with _amount equal to 0. Line 5 is respected and then pending variable is estimated. Since the user.rewardDebt equal to 0, the equation becomes pending = user.amountWithBonus.mul(pool.accRadsPerShare).div(1e12). Both pool.accRadsPerShare and user.amountWithBonus are positive number, so the product pending larger than 0 as well. Since the statement at line 10 is not respected, the code can't adjust the user.amountWithBonus variable to indicate the user claims the reward.

```
1 function withdraw(uint256 _pid, uint256 _amount) external validatePool(_pid) {
2   PoolInfo storage pool = poolInfo[_pid];
3   UserInfo storage user = userInfo[_pid][msg.sender];
4   require(user.amount >= _amount, "withdraw: not good");
5
6   updatePool(_pid);
7   uint256 pending = user.amountWithBonus.mul(pool.accRadsPerShare).div(1e12).sub(user.
rewardDebt);
8   if(pending > 0) {
9     if(pool.isSRadsRewards){
10       safeSRadsTransfer(msg.sender, pending);
11     }
12     else{
13       safeRadsTransfer(msg.sender, pending);
14     }
15   }
16   if(_amount > 0) {
17     user.amount = user.amount.sub(_amount);
```

```
18         uint256 _bonusAmount = _amount.mul(userBonus(_pid, msg.sender).add(10000)).div
        (10000);
19         user.amountWithBonus = user.amountWithBonus.sub(_bonusAmount);
20         pool.lpToken.safeTransfer(address(msg.sender), _amount);
21         pool.lpSupply = pool.lpSupply.sub(_bonusAmount);
22     }
23     user.rewardDebt = user.amountWithBonus.mul(pool.accRadsPerShare).div(1e12);
24     emit Withdraw(msg.sender, _pid, _amount);
25 }
```

Listing 3.16: Deposit Function

The `user.amountWithBonus` increases every time the attacker starts from the step 1. This enables the attacker to drain more and more tokens in the process. Checking the transaction on BSCscan, it is shown how many times the attacker replicated this methodology.

3.7.2 Properties

The logic of the smart contract vulnerability is contained in the procedure of estimation of users' rewards.

The malicious sequence of functions involves the call of `deposit`, `emergencyWithdraw` and `withdraw`. Therefore, the attacker could get back the same amount of deposited tokens, but with a higher amount bonus for the reward. The parameter `amountWithBonus` of user struct, which keeps track of the amount and the bonus, just increases even if the user receives the reward and it is withdrawing.

For the detection of the vulnerability, we specify the property as a postcondition of the function `withdraw`. Considering the caller, who receives a reward because of the bonus, the parameter of the struct `amountWithBonus` has to decrease.

3.8 Reentering the Reentrancy Bug: Disclosing BurgerSwap's Vulnerability

BurgerSwap is an automated Market Maker service on Binance Smart Chain (BSC). At time of the disclosure of the vulnerability, there was around \$13K worth of Ether at immediate risk. The vulnerability was discovered by <https://zengo.com> ZenGo team and it was presented by **BurgerSwap**.

3.8.1 BurgerSwap

BurgerSwap is a Binance Smart Chain fork of Uniswap, Automated Market Maker (AMM) service operating on Ethereum. Trading and listing Specialized BEP-20 tokens among standard swapping options are available on this platform. To mint such tokens, users can use BurgerSwap's "bridge" contract on Ethereum. Ethereum-BSC "bridge" contract on Ethereum was the main target of the attack.

Brige is a combination of 2 smart contracts deployed on different chains. It allows cross-chain transfers of value. Ether deposited into the contract on the main net will provide a balance denominated in ERC-20 tokens on the sidechain. While ERC-20 tokens deposited back into the contract on the sidechain can free up Ether on main net. One example could be locking Ether, which is converted via the contract to WETH (Wrapped Ether, an ERC-20 token pegged to Ether), and then the same wallet locking ETH can be credited with bWETH on BSC.

3.8.2 The vulnerability

The issue deals with the function `withdrawFromBSC`, ???. First of all, it checks some conditions and then it proceeds to transfer the amount to the message sender. The order of the actions is:

1. It verifies `executeMap[_paybackId]` is false;
2. It checks `_signature` is a valid signature on `_paybackId`, `_token`, `msg.sender`, and `_amount`.
3. It calls `TransferHelper.safeTransferETH(msg.sender, _amount)`.
4. It sets `executeMap[_paybackId]` to true.

The issue is the interaction with the sender's address (step 3) happens before the internal effect (step 4): reentrancy is feasible.

```
1 library TransferHelper {
2     function safeApprove(address token, address to, uint value) internal {
3         // bytes4(keccak256(bytes('approve(address,uint256)')));
4         (bool success, bytes memory data) = token.call(abi.encodeWithSelector(0x095ea7b3, to
5         , value));
6         require(success && (data.length == 0 || abi.decode(data, (bool))), 'TransferHelper:
7         APPROVE_FAILED');
8     }
9
10    function safeTransfer(address token, address to, uint value) internal {
11        // bytes4(keccak256(bytes('transfer(address,uint256)')));
12        (bool success, bytes memory data) = token.call(abi.encodeWithSelector(0xa9059cbb, to
13        , value));
14        require(success && (data.length == 0 || abi.decode(data, (bool))), 'TransferHelper:
15        TRANSFER_FAILED');
16    }
17
18    function safeTransferFrom(address token, address from, address to, uint value) internal
19    {
20        // bytes4(keccak256(bytes('transferFrom(address,address,uint256)')));
21        (bool success, bytes memory data) = token.call(abi.encodeWithSelector(0x23b872dd,
22        from, to, value));
23        require(success && (data.length == 0 || abi.decode(data, (bool))), 'TransferHelper:
24        TRANSFER_FROM_FAILED');
```

```

20     function safeTransferETH(address to, uint value) internal {
21         (bool success,) = to.call{value:value}(new bytes(0));
22         require(success, 'TransferHelper: ETH_TRANSFER_FAILED');
23     }
24 }
25
26 contract ETHBurgerTransit {
27     ...
28     function withdrawFromBSC(bytes calldata _signature, bytes32 _paybackId, address _token,
29         uint _amount) external payable {
30         require(executedMap[_paybackId] == false, "ALREADY_EXECUTED");
31
32         require(_amount > 0, "NOTHING_TO_WITHDRAW");
33         require(msg.value == developFee, "INSUFFICIENT_VALUE");
34
35         bytes32 message = keccak256(abi.encodePacked(_paybackId, _token, msg.sender, _amount
36             ));
37         require(_verify(message, _signature), "INVALID_SIGNATURE");
38
39         if(_token == WETH) {
40             IWETH(WETH).withdraw(_amount);
41             TransferHelper.safeTransferETH(msg.sender, _amount);
42         } else {
43             TransferHelper.safeTransfer(_token, msg.sender, _amount);
44         }
45         totalFee = totalFee.add(developFee);
46
47         executedMap[_paybackId] = true;
48
49         emit Withdraw(_paybackId, msg.sender, _token, _amount);
50     }
51     ...
52 }

```

Listing 3.17: BugerSwap Bridge Contract

Following the execution of the code, the bug is found in the `safeTransferETH` function, line 23, contained in `TransferHelper` library. The expression `to.callvalue:value(new bytes(0))` is actually a call to the sender of the message, which can be an arbitrary smart contract. The malicious contract can implement a fallback function. By the time it receives the ether, the fallback function is triggered and `withdrawFromBSC` is run again, but without updating `executeMap[_paybackId]`. Since it is not set to true, the code repeat the same sequence of operation. Repeating this process within the same transaction, the attacker will drain the vulnerable contract's WETH holdings and credit.

3.8.3 Prorties

The keyword of this exploit is reentrancy.

The smart contracts adopt an access control strategy at the beginning of the function, checking the signature of the user. The attacker forked `MetaMask`, a crypto-wallet, for allowing a smart contract to access the vulnerable function. However, it is out of our interest and we focus on the vulnerability in the code.

The vulnerable function is `safeTransferETH`, responsible of sending ETH to the user. This is called by the one which manages the withdrawal of funds from the contract by the user. The malicious contract, thanks a fallback, can call multiple times the same function and withdrawing more money than it could.

With the postcondition, we check the balance of WETH (wrapped ETH) before and after the function, stating the difference should be the parameter amount of the function.

3.9 Infinite minting of NFTs

Introduction

3.9.1 DirtyDogs NFT

The logic of contract is -> presale sold like ticket -> end of presale, claim ticket as NFT -> mint until a limit

3.9.2 The exploit

DirtyDogs NFT contract has a typical example of reentrancy. The attacker exploited the function `claimDogs()`, shown in ?? . Firstly, the malicious wallet bought a ticket for having the right of receiving a NFT, calling the function `claimDogs()`. It basically loops on the number of ticket the sender has, and it calls the function `_safeMint` for creating the NFTs and sending them to the caller. The bug stays in line 31, because `totalClaimed[_msgSender()]` is updated at the end of the loop. It is the variable which keeps track of the number of tickets owned by the caller.

The fundamental step of the exploit was the caller of the function: a smart contract. It implemented a callback function: main trigger for reentrancy attacks. Within the same transaction, it gets the opportunity to execute the some code multiple times. When the smart contract receives an NFT, the fallback function is triggered and the `claimDogs()` function is called again. As result, the attacker could call again the function for minting, but without updating the variable which counts the number of ticker per address. The exploit produced 45 NFTs, because the fallback has the risk of revert, there is a limit of times to be called.

```
1 contract ERC721 is Context, ERC165, IERC721, IERC721Metadata, IERC721Enumerable {
2     ...
3     function _mint(address to, uint256 tokenId) internal virtual {
4         require(to != address(0), "ERC721: mint to the zero address");
5         require(!_exists(tokenId), "ERC721: token already minted");
6
7         _beforeTokenTransfer(address(0), to, tokenId);
8
9         _holderTokens[to].add(tokenId);
10
11         _tokenOwners.set(tokenId, to);
12
13         emit Transfer(address(0), to, tokenId);
```

```

14     }
15
16     ...
17 }
18
19     ...
20
21 contract DirtyDogs is ERC721, Ownable {
22     ...
23     function claimDogs() external {
24         uint256 numbersOfTickets = getUserClaimableTicketCount(_msgSender());
25
26         for(uint256 i = 0; i < numbersOfTickets; i++) {
27             uint256 mintIndex = totalSupply();
28             _safeMint(_msgSender(), mintIndex);
29         }
30
31         totalClaimed[_msgSender()] = numbersOfTickets.add(totalClaimed[_msgSender()]);
32     }
33
34     function getUserClaimableTicketCount(address user) public view returns (uint256) {
35         return presaleNumOfUser[user].add(publicNumOfUser[user]).sub(totalClaimed[user]);
36     }
37     ...
38 }

```

Listing 3.18: DirtyDogs NFT contract

3.9.3 Properties

The attacker could exploit the smart contract due to a vulnerability in the function `claimDogs`, which includes even a bad implementation of the ERC721 standard. It can be classified as a case of reentrancy.

The NFTs are not directly sold, but a ticket instead is provided to the users, who would convert it. The function `claimDogs` is in charge to verify the tickets and generate the NFTs. Because of the reentrancy, the attacker produced 45 NFTs with a single ticket.

The property is a postcondition, which states that the amount of NFTs produced should be equal to the number of tickets of the user.

4 Analysis Tools

This chapter is addressed to present the analysis tools involved in the experiments. Each one has a dedicated section, which gives a general overview of its architecture, the implemented methodologies and the limitations. We collected the informations from their presenting paper or the provided github documentations.

Our work deals with the practical aspect of the tools, their installation and their behavior in a real world case, consequently, we decided to stress this practical aspect for defining their taxonomy. We draw it, classifying the tools based on their running mode. ?? depicts our classification, involving three different subgroups:

- tools without specification;
- tools with specification;
- tools with multiple running mode.

The first one covers the ones which can be run with the source code (or bytecodes) of the smart contract without any additional data. Their strategies can detect a predefined group of vulnerabilities. On the other hand the second group involves the ones which require the user provides specifications for running. The specification can involve a rule to be broken or a property which can be proved or unproved. In our taxonomy, the meaning of "multiple mode running" implies the tools which can be run with and without provided specifications.

The choice of these eight analysis tools is based on our literature research, which encompasses surveys and comparisons of those. Another discriminant aspect of the choice is the maintenance of the tools. We decided to discard the ones which have not been updated for more than two years. We decided to consider just open-source tools because their architecture descriptions and documentation are available. We made an exception in Certora, which is not open source. It has the premium version and the free one, we adopted the second one and the free provided documentation online.

Table 4.1: Tools Classification

Tools with Specification	Tools without Specification	Tools with Multiple Running Mode
SolcVerify	Mythril	Manticore
Celestial	Slither	SmarTest
Echidna		
Certora		

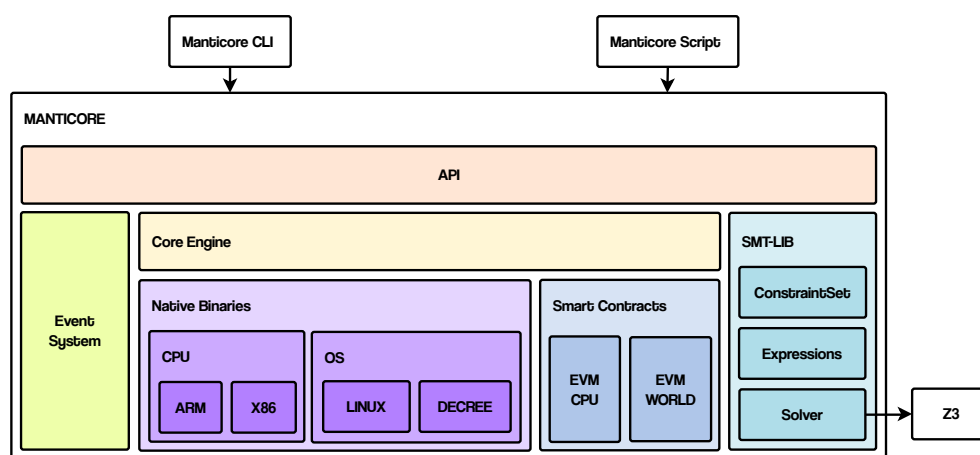


Figure 4.1: Manticore Architecture

4.2.1 Certora Prover

Certora Prover is one of the most well-known and used tools for formal verification of Solidity smart contracts. Any computer programme that may be compiled using EVM can undergo Certora Prover verifications.

Since the tool is not open source, we draw the information regarding the tool from the **Certora Documentation**

It is provided as Software as a Service, a cloud technique, so it is not possible to install the complete tool. A user can interact with it on its website, otherwise, a command-line interface can be downloaded, which interacts with the server by remote.

The user for the verification has to provide the Solidity file and the specification one, which contains the logic formulas for verification conditions; these are proven by an SMT (satisfiability modulo theories) solver. The result is the proof of the rule, otherwise, if the rule is disproved, the solver also provides a concrete test case demonstrating the violation.

A user can state the specifications as invariants otherwise as functions, called rules. The invariants contain just a boolean formula. On the other hand, the rules contain the properties to be proved and these are written like functions. For coding these one, a similar language to Solidity is used. A rule can contain requires at the beginning, for expressing a condition to be realised before the running of it. Functions from the smart contract can be called. A requirement is that each rule has to conclude with assert, containing a boolean condition.

?? covers an example of the specification of the function `transferFrom()` of a smart contract which implements a token. The function cares about the transfer of tokens from one account to the other one. This rule checks that the balances of the users are updated correctly.

```

1
2 rule transferFromCorrect(address from, address to, uint256 amount) {
3   env e;
4   require e.msg.value == 0;
5   uint256 fromBalanceBefore = balanceOf(from);
6   uint256 toBalanceBefore = balanceOf(to);

```




Figure 4.2: Celestial Architecture

```

7  uint256 allowanceBefore = allowance(from, e.msg.sender);
8  require fromBalanceBefore + toBalanceBefore <= max_uint256;
9
10 transferFrom(e, from, to, amount);
11
12 assert from != to =>
13     balanceOf(from) == fromBalanceBefore - amount &&
14     balanceOf(to) == toBalanceBefore + amount &&
15     allowance(from, e.msg.sender) == allowanceBefore - amount;
16 }

```

Listing 4.3: Certora example specifications

4.2.2 Celestial

This subsection is addressed to introduce Celestial, an analysis tool for Solidity Ethereum-based smart contracts developed by the research team of Microsoft India. ?? shows the its architecture.

The developes provide functional requirements for formally verifying their specifications. The input file is labelled It gives programmers the ability to create functional requirements for their contracts. The input file is labelled as ".cel", it is the solidity file, with the added specification expressed in notes. When the grammar is checked, the contract and the specifications are translated in F* for having the verified verdict.

?? shows an example of input file. The invariants are expressed in a sort of functions. At the beginning of a function, the specification can be expressed, regarding precondition, postcondition and so on. One of these can involve the keyword modifies, which specifies the variable that can be modified in the function, or tx_reverts, which states the possible condition that a function can revert. The Solidity implementation of the function is kept.

```

1  contract SimpleMarketplace {
2      // contract fields
3      invariant balanceAndSellerCredits {
4          balance >= totalCredits &&

```

```
5         totalCredits == sum_mapping ( sellerCredits )
6     }
7     //function
8     function buy ( address itemId ) public
9         modifies [ sellerCredits , totalCredits , itemsToSell ,
10             log ]
11         tx_reverts !( itemId in itemsToSell ) || value != itemsToSell [ itemId ].price
12         || value + totalCredits > uint_max
13         post (!( itemId in itemsToSell ) && sellerCredits [ seller ] == old (
14             sellerCredits ) [seller => sellerCredits [ seller ] + value ]
15         && log == ( eItemSold , sender , itemId ) :: old ( log ) )
16     { // implementation of the buy function }
```

Listing 4.4: Celestial example specifications

F* is a fully dependent type system proof helper and programs verification. The authors gave the same reasons for involving F* for the formal proof in a blockchain context. First, it offers SMT-based automation, which is sufficient for the completely automated verification of real-world smart contracts. Second, F* enables the developers to work in a customised state and exception effect mimicking the blockchain semantics since it supports user-defined effects. Finally, even though we only use its first-order subset with quantifiers and arithmetic, F* permits expressive higher-order specifications.

Celestial process involves 2 steps: the translation of the specification and the verification of F* start. The first one involves a python script, on the other hand the second one entails the intrallation of F* engine. The output covers the response of the verification and a generated solidity file, which represents the smart contract without the specifications notes.

Limitations The authors explained their tool implementation focused on the Solidity constructs used in their case studies, therefore it does not cover some Solidity cases.

Delegatecall, embedded assembly It does not take into account syntactic elements like inheritance, abstract contracts, tuple types, delegatecall and embedded assembly

Most of these only offer syntactic sugar, which CELESTIAL's future iterations should find simple to support. Arrays and structs cannot presently be passed as parameters to functions in our implementation.

Loops are allowed in the smart contracts, however the tool does not support loops invariants. When external contracts are called, reentrant behaviour can result, in which the external contract contacts the caller back. Reasoning about reentrant actions is frequently counterintuitive. Celestial forbids these actions, this property is called "external callback freedom" (ECF). It states that every callback execution in a contract is equivalent to some activity without reentrancy. So Celestial assumes that there is no callback during the external call. Programmers can use the tool to create and support the specifications of their own contracts without making any assumptions about the behaviour of external contracts.




Figure 4.3: Echidna architecture

4.2.3 Echidna

Echidna is an open-source smart contract fuzzer, developed by **Echidna**, which makes it easy to automatically generate tests to detect violations in assertions and custom properties. Rather than relying on a fixed set of pre-defined bug oracles to detect vulnerabilities during fuzzing campaigns, Echidna supports three types of proper- ties:

- user-defined properties (for property-based testing);
- assertion checking;
- gas use estimation.

Figure ?? depicts the Echidna architecture as a two-step process: pre-processing and fuzzing. The tool starts with a collection of contracts that have been supplied, as well as attributes that have been integrated into one of the contracts. Echidna uses Slither , smart contract static analysis framework presenet in ??, to build and analyse the contracts in order to find relevant constants and functions that directly handle Ether (ETH). The fuzzing effort begins in the second stage. Using the application binary interface (ABI) given by the contract, significant constants stated in the contract, and any previously gathered sets of transactions from the corpus, this iterative procedure creates random transactions. When a property violation is detected, a counterexample is created to indicate the smallest and most basic sequence of operations that caused the failure.

The code ?? provides an example of invariant in Echidna context. The Solidity contract contains a vulnerability a the backdoor function. The output of the terminal is presented in Listing ??: the attacker. For breaking the property, can call in order the tunctions airdrops() and backdoor()

```

1 contract Token{
2     mapping(address => uint) public balances;
3     function airdrop() public{
4         balances[msg.sender] = 1000;
5     }
6     function consume() public{
7         require(balances[msg.sender]>0);
8         balances[msg.sender] -= 1;
9     }
10    function backdoor() public{
11        balances[msg.sender] += 1;
12    }
13    function echidna_balance_under_1000() public view returns(bool){

```

```
14     return balances[msg.sender] <= 1000;
15 }
16 }
```

Listing 4.5: Solidity smart contract implementing a vulnerable Token and an Echidna invariant function.

The tool can be even used to test assertions. The aim is equivalent of the invariant testing methodology, but in this case properties are expressed using the Solidity annotation of assertion.

```
1 $ echidna-test testtoken.sol --contract TestToken
2 ...
3 echidna_balance_under_1000: failed!
4 Call sequence, shrinking (1205/5000):
5 airdrop()
6 backdoor()
7
8 ...
```

Listing 4.6: Tool’s result after the execution of the precious code.

4.2.4 Solc-Verify

SolcVerify present solc-verify, a source-level verification tool for Ethereum smart contracts. It takes smart contracts written in Solidity and discharges verification conditions using modular program analysis. It is built on top of the Solidity compiler, so it reasons at the level of the contract source code. Because of that, Solc-verify is able to reason about high-level contract attributes while accurately modeling low-level language semantics.

Solc-verify is implemented as an extension to the Solidity compiler. It accepts a collection of Solidity contracts, including specification annotations, and uses the Boogie verifier and SMT solvers to discharge verification conditions.

As **SolcVerify_2** explain, Solc-verify translates the annotated contracts to the Boogie Intermediate Verification Language (IVL). The key idea of the translation is to encode state variables as global heaps and functions as procedures. Solc-verify relies on the Boogie verifier to perform modular verification by discharging verification conditions to SMT solvers. The verification conditions encode the function body while assuming the preconditions, and then check if postconditions hold. In this process, function calls are replaced by their specification and loops by their invariants (modularity). Finally, the results are back-annotated to the Solidity source.

?? present an example of annotation, which states that the contract will ensure that the sum of individual balances is equal to the total balance in the bank.

```
1 pragma solidity >=0.7.0;
2
3 /**
4  * @notice invariant __verifier_sum_uint(balances) <= address(this).balance
5  */
6 contract SimpleBank {
7     mapping(address=>uint) balances;
```



```

8
9  function deposit() public payable {
10      balances[msg.sender] += msg.value;
11  }
12
13  function withdraw(uint256 amount) public {
14      require(balances[msg.sender] > amount);
15      bool ok;
16      (ok, ) = msg.sender.call{value: amount}(""); // Reentrancy attack
17      if (!ok) revert();
18      balances[msg.sender] -= amount;
19  }
20 }

```

Listing 4.7: An example Solidity smart contract implementing a simple bank with SolcVerify annotations.

SolcVerify_3 on GitHub repository, present the specification annotations. Those must be included in special documentation comments (`///` or `/** */`) and must start with the special doctag `@notice`. They must be side-effect free Solidity expressions (with some verifier specific extensions) and can refer to variables within the scope of the annotated element. Functions cannot be called in the annotations, except for getters. The currently available annotations are listed below.

- Function pre/postconditions can be attached to functions. Preconditions are assumed before executing the function and postconditions are checked (asserted) in the end. The expression can refer to variables in the scope of the function. The postcondition can also refer to the return value if it is named.
- Contract level invariants can be attached to contracts. They are included as both a pre- and a postcondition for each public function. The expression can refer to state variables in the contract (and its balance).
- Loop invariants can be attached to for and while loops. The expression can refer to variables in scope of the loop, including the loop counter.
- Modification specifiers can be attached to functions. The target can be a (1) state variable, including index and member accesses or (2) a balance of an address in scope. Note however, that balance changes due to gas cost or miner rewards are currently not modeled.
- Event data specification can be attached to events that should be emitted when certain data changes. Events can declare the state variable(s) they track for changes, or in other words, the variables for which the event should be emitted on a change.

4.3 Tools without specification

We introduce the second main category of our taxonomy: tools without Specifications.

These do not need any specifications and they based their analyses on preselected strategies. They can be considered "plug & play", indeed they work with code (or even bytecode) of the smart contracts, without the developer should code anything for letting them run. Most of the tools that we found during the literary research propose just a warning, depicting the vulnerable lines of code and classifying the vulnerability following their taxonomy or open-source ones.

4.3.1 Slither

Slither is described by **Slither** as an open-source static analysis framework. It uses its own intermediate representation, SlithIR, which was created to simplify static analysis of Solidity code. Concolic analysis, taint analysis, and control flow checking are involved for detecting a variety of security vulnerabilities. It is designed to provide granular information about smart contract code and the flexibility necessary to support many applications.

It is mainly used for:

- Automated vulnerability detection: a large variety of smart contract bugs can be detected without user inter- vention.
- Automated optimization detection: Slither detects code optimizations that the compiler misses.
- Code understanding: printers summarize and display contracts' information to aid in the study of the codebase.
- Assisted code review: through its API, a user can interact with Slither.

Slither implements more than twenty bug detectors, regarding reentrancy, Uninitialized variables, Shadowing and many other. The tool allows the developers to integrate more detectors, therefore it extends Slither's capabilities to detect more advanced bugs.

SlitherGitHub is written in python 3 and it is published on GitHub. During the installation, I did not find any particular issues.

4.3.2 Mythril

Mythril is a security analysis tool for Ethereum smart contracts. It was introduced by **Mythril**.

The tool relies on concolic analysis, taint analysis and control flow checking of the EVM bytecode to prune the search space and to look for values that allow exploiting vulnerabilities in the smart contract. It is targeted at finding common vulnerabilities, and is not able to discover issues in the business logic of an application. **SWCRegistry**'s taxonomy of vulnerabilities is used by Mythril for classify them. Listig ?? illustrates an example of output of Mythril analysis. At the secocond line, there is the reference to the vulnerability classified by SWC Registry with the ID of 110 (Assert Violation).

```
1 ==== Exception State ====
2 SWC ID: 110
3 Severity: Medium
```

```
4 Contract: Token
5 Function name: transferArray(address[],uint256[])
6 PC address: 4385
7 Estimated Gas Usage: 944 - 6585
8 An assertion violation was triggered.
9 It is possible to trigger an assertion violation. Note that Solidity assert() statements
  should only be used to check invariants. Review the transaction trace generated for this
  issue and either make sure your program logic is correct, or use require() instead of
  assert() if your goal is to constrain user inputs or enforce preconditions. Remember to
  validate inputs from both callers (for instance, via passed arguments) and callees (for
  instance, via return values).
10 -----
11 In file: test.sol:309
12
13 function transferArray(address[] tos, uint256[] values) public returns (bool) {
14     for (uint8 i = 0; i < tos.length; i++) {
15         require(transfer(tos[i], values[i]));
16     }
17
18     return true;
19 }
20
21 -----
```

Listing 4.8: Example of the output of Mythril Analysis.

5 Outcome for individual tools

This chapter is moved in the next one, I keep now bc it contains the notes

Fomal specifications, outcomes, time of running, no comparison, single description, outliers (performe very well or very baddly), difficulties in the installation and running

Detalais regarding what I did for running the test, so for each tool I describe the setting for each scripts.

5.0.1 SmarTest

In some cases it takes to long so I added a timer of 320 with the setting no reentrancy by definition set up in exploit mode + assertion mode About Aku it gives half of the result, it states that all the user has to be refunded, but the

5.0.2 Slither

AKU: it detects the presence of arbitrary call but it doesn't detect that can get stucked Correct detection of all the reentrancy Cover: warning about the comparison based on blocktimestamp

it has no capability for guess the other vulenrabilities

5.0.3 Mythril

AkU: nada Correct detection of all the reentrancy Cover: warning about the comparison based on blocktimestamp

5.0.4 SolcVerify

Notes:

- inheritance allowed
- it worked well with the codes (Solidity Version 7)
- the analysed files are still in solidity
- grammar language very intuitive
- it allwos even not flat contract

It works better than Celestial

5.0.5 Celestial

Notes:

- no inheritance
- no struct and no array as parameter in function
- total different language
- more grammar for the language
- A lot of grammar
- no address(this), it gives error in fstar
- problem with internal functions , fstar could not find the identifier (BZX)
- not so clear which property is violeted
- no loop, neither invariant neither written

Uranium: false positive

Spartan: it works good

bZX: about properties it worked properly, problem with internal functions

BurgerSwap & DirtyDogs & SurgeProcol: no reentrancy bc Calling external contracts are avoided. Contract Local Reasoning: Calling external contracts can lead to reentrant behavior where the external contract calls back into the caller, which is often difficult to reason about. CELESTIAL disallows such behaviors by checking for external callback freedom (ECF) [28], [42] which states that every contract execution that contains a reentrant callback is equivalent to some behavior with no reentrancy. When this property holds, it is sufficient to reason about non-reentrant

Aku: it didn't work with the loop so rewrite the contract, but it worked

Cover Protocol: celestial doesn't recognize memory and storage keyword

5.0.6 Manticore

BurgerSwap & DirtyDogs & SurgeProcol: no reentrancy, bc it gives always a fake positive
You can test the Reentrancy just with the tool without specification, running Manticore as Scanner

Spartan: worked well Uranium: fake positive

Aku: Worked well , it says that the rule about claimProject is broken just if everybid is

1

Cover: doesn't work

5.0.7 Echidna

Spartan: it worked correctly and it gave even the transaction order, tested with assertions and with test Function

Aku: worked well , it found the correct sequence of transaction for block the contract

Uranium: it worked properly, it gave the list of actions to do

Cover Protocol: worked good

BZX: verified with the assertion but it worked properly

5.0.8 Manticore

5.0.9 Certora

The time of execution is given by the application

5.0.10 Tool without Specification

Mythril and Slither Slither super fast Mythril in case it was too slow we put a timer of 240 seconds A timer was added in python modifying the main caller file

6 Evaluation

We collected the outcomes of the analyses of each tool. In this chapter, we give a general overview of what we obtained, firstly, describing the result per tool, then comparing all of those.

6.1 Outcomes per Tool

This section encompasses the result for each tool. We show per each tool the outcomes of each running per attack.

Manticore The provided guide on Github ([**ManticoreGitHub**]) gives a detailed guide for Manticore installation. Since it is written in python, we used a virtual environment and we counted around nine libraries for dependences. The installation involved basically one command, since it was managed by "pip", python packet manager. We adopted its default running mode for the attacks involving reentrancy and "Manticore-verifier" running mode for the others.

The table ?? shows the outcomes of the analyses per attack; the symbol "-" means that in that case we did not use the "Manticore-verifier" running mode, but the default one, which it does not need a specification file.

SmarTest SmarTest is built on top of VeriSmart tool, so it can be seen like a plug in of these one. Indeed, we run SmarTest as an option of VeriSmart, as the Github guide explains ([**SmarTestGitHub**]). The tool is built with OCalm, a program language, so we used for

Table 6.1: Manticore results

Attacks	Line of Code for Specifications	Time of Execution (in seconds)
Aku	6	235
Cover	5	245
BZX	4	228
Spartan	3	239
Uranium	3	250
XSURGE	–	208
BurgerSwap	–	205
DirtyDogs	–	203

Table 6.2: SmarTest results

Attacks	Line of Code for Specifications	Time of Execution
Aku	3	320
Cover	2	310
BZX	4	320
Spartan	2	320
Uranium	4	320
XSURGE	–	–
BurgerSwap	–	–
DirtyDogs	–	–

Table 6.3: Celestial results

Attacks	Line of Code for Specifications	Time of Execution
Aku	22	4
Cover	–	–
BZX	15	3
Spartan	29	5
Uranium	18	5
XSURGE	–	–
BurgerSwap	–	–
DirtyDogs	–	–

the installation "opam", which is a source-based package manager for it. An important dependency is Z3, its satisfiability modulo theories (SMT) solver. Solc, the compiler for solidity, is required.

Since it has no the detector for reentrancy, as the table ?? shows, the attacks involving it were discarded. We adopted just the "assertion" running mode, for obtaining valuable results. We fixed a running time threshold of 320 seconds.

Celestial

Echidna

Certora

Table 6.4: Echidna results

Attacks	Line of Code for Specifications	Time of Execution
Aku	7	22
Cover	5	7
BZX	3	33
Spartan	3	17
Uranium	3	24
XSURGE	–	–
BurgerSwap	–	–
DirtyDogs	–	–

Table 6.5: Certora results; the time is provided by the sas application

Attacks	Line of Code for Specifications	Time of Execution
Aku	52	14
Cover	31	21
BZX	25	18
Spartan	20	25
Uranium	42	27
XSURGE	–	–
BurgerSwap	–	–
DirtyDogs	–	–

Table 6.6: SolcVerify results

Attacks	Line of Code for Specifications	Time of Execution (seconds)
Aku	9	4
Cover	13	5
BZX	17	9
Spartan	25	17
Uranium	23	9
XSURGE	20	10
BurgerSwap	11	10
DirtyDogs	30	14

SolcVerify

Mythril

Slither

6.2 Comparison of Analyses

The first step was the collection of all the data from the analyses with the vulnerable smart contracts as objective. In this part, we provide a technical comparison of what we found.

?? has as rows the tools and as columns the real-world exploits. It shows which tool was able to scan the specified vulnerability involved in the attacks. The caption states the symbol we have used for the classification:

- ✓ states that the vulnerability was scanned;
- ✗ means the tool was not able to detect it;
- – stands for "discarded", the attack was not considered for architectural reason of the tool.

We can state that the tools with specifications could not detect reentrancy in most of the cases. We classify that case as discarded. SolcVerify is the only one which could provide a result in all cases, moreover, it was the only tool which could detect correctly all the vulnerabilities we provided.

Celestial could not analyse BZX exploit, because it involved the keywords "storage" and "memory", which cannot be scanned by the tool.

The tools without specifications gave warning in most cases, which can be used as a hint for scanning the vulnerabilities.

Slither could not detect the exploits involving the logic of the program, moreover, it gives multiple warnings, which the developer took into account. An example is BZX, which

Table 6.7: Analyses Outcomes per Attack: ✓: Found vulenrability, ✗: Not found vulnerability, -: Discarded

Tools	Aku	Cover	BZX	Spartan	Uranium	XSURGE	BurgerSwap	DirtyDogs
Manticore	✗	✗	✓	✓	✗	✓	✓	✓
SmartTest	✓	✗	✓	✗	✓	-	-	-
Celestial	✓	-	✓	✓	✓	-	-	-
Echidna	✓	✓	✓	✓	✓	-	-	-
Certora	✓	✓	✓	✓	✓	-	-	-
SolcVerify	✓	✓	✓	✓	✓	✓	✓	✓
Slither	✗	✗	✗	✗	✗	✓	✓	✓
Mythril	✗	✗	✗	✗	✗	✓	✓	✓

Table 6.8: Analyses Outcomes: LoF: List of functions, LoU: List of unproved tests, W: Warnings

Tools	Constructive output	Avg lines of code for test	Avg time (in seconds)
Manticore	LoF, W	4	239,5
SmartTest	LoF, W	2,5	318
Celestial	LoU	21	4
Echidna	LoF	4	20,5
Certora	LoF	34	21
SolcVerify	LoU	18,5	10
Mythril	LoF, W	-	221
Slither	W	-	3,5

involves the estimation time comparison based on timestamp, but for the tool, a comperison timestamp based is always classified as dangerous.

Exploits with no reentrancy We want to focus

Tools behavior

6.3 Weaknesses and Strenghts

We present the

Table 6.9: User Expirience

Tools	Running modes	Extenral Dependences	OS
Manticore	2	2	Linux, OS X
SmartTest	2	3	Linux, OS X, Windows
Celestial	1	3	Linux, OS X
Echidna	2	2	Linux, OS X, Windows
Certora	1	–	–
SolcVerify	1	4	Linux, OS X
Mythril	1	2	Linux, OS X
Slither	1	2	Linux, OS X, Windows

Table 6.10: Weaknesses & Drawbacks

Tools	Weaknesses & Drawbacks
Manticore	Reentrancy is not dectected by properties property based execution, very slow
SmartTest	Reentrancy is not dectected, the analyses are slow
Celestial	External calls are not considered, keywords storage and memory are not recognized
Echidna	Reentrancy is not dectected, no assestion mode for solidity 8
Certora	Reentrancy is not dectected, not open-sources
SolcVerify	It just gives warning, it does not provide a list of transaction for breaking the given property
Mythril	Just flat contracts are allowed
Slither	Scan is based on the grammar, great amount of false negative

Table 6.11: Strenghts

Tools	Strenghts
Manticore	One of the mode cover the properties breaking and the scanner one covers the reentrancy
SmartTest	It allows to set the specific vulenrability to look for
Celestial	Possibility to use different version of F*
Echidna	Possibility to run in multiple modes with different grammar (tests or assertions breaking)
Certora	Implements the library of Openzeppelin, SAS no installation needed
SolcVerify	Intuitive specification language based on Annotations, it detects reentrancy
Mythril	It deos not need specification, but still provide list of functions for breaking detected vulnerab
Slither	Easiest installation, fastest tool that we used

7 Discussion

In this chapter we discuss the obtained result and we give our personal point of view.

7.1 Tools with Specifications

The considered tools cover three main strategies for security analysis: fuzzing, symbolic execution and formal verification.

The following paragraphs deal with comparison within this subgroup we want to stress for having a deepen view of the comparison and a last sub chapter for a general overview.

Formal verification Formal verification is a very powerful security approach, intending to prove or unproved the given specification. This perfectly fits with our research goals, such as the detection of bugs or vulnerabilities in our real-world cases.

We involved three different tools, implementing this approach, for our purpose. Regarding the results, it is clear that these were executed as fast as or even faster than the other ones. Certora is the only one which provides a complete list of functions for breaking the rules, rather than just a warning. On the other hand, SolcVerify could detect the vulnerabilities involving external call functions, indeed reentrancy. A powerful aspects of this tool is its possibility to express loop invariants, the other ones do not allow it. Considering the grammar for expressing the specifications, SolcVerify is the one which needs the least amount of lines of code, indeed it involves a notification language; I found it very intuitive and fast to write down the specifications.

Celestial architecture encompasses two steps: the translation from celestial file to f* and then its verification. The python script converts the ".cel" in f*, used for the proof or unproof. The provided file included the smart contract's source code plus the expressed specification. These are statements placed at the beginning of a function, otherwise, it is possible to create a sort of function, containing boolean formula, which is called by different specifications function with different parameters, it is useful for expressing the same specification for different purpose. I consider it the one with more limitations, regarding solidity grammar and reentrancy, because it could not detect the reentrancy vulnerabilities and the Cover protocol attack, because the keywords "storage" and "memory" are rejected.

Certora is the only tool which is not open-source, for our purpose we adopted its free version. Its specification language is described by its developers' group as "rule-based". It differs from the other two tools under this aspect, because this way gives more elasticity to the user and defines more specific cases. The rule is composed of some function calls and it concludes with an assertion or more. The user is allowed to test a specific case, using "require" and the possibility to set up a proper environment. The preconditions, in this case, are expressed using the Solidity keyword "require" in the rule.

One of its strengths is the possibility to define the specifications we want to prove, without necessarily defining all the specifics for the rest of the functions. On the other hand, with Celestial and SolcVerify, we provided the specifications for all the functions for letting the tools work properly. Those could not prove the given properties, without the specifications for all the code.

Different strategies, similar grammar We considered two tools with similar grammar but implemented different analysis approaches: Manticore and Echidna. In both of the cases, we provided functions containing a boolean formula, which the tools try to break. From the results, we noticed that Echidna run faster and it worked for all the cases, but Manticore could cover the reentrancy vulnerabilities thanks to its changing architecture.

We encountered a common aspect between the grammar definition of the specification between the two tools Echidna and SmartTest regarding their "assertion" mode. Both of those require the user to write assertions and then these try to verify it or return the list of functions for breaking the rule. From our results, it is clear Echidna could obtain higher number of positive outcomes and in less time rather than SmartTest.

7.2 Customized and Non-specific Analyses

The objectives of our analyses were smart contracts involved in real-world exploits. The attackers exploited a specific bug or lack of security in the logic of programs.

The specifications allowed the user to express the requirements of the program. These provide a customized analysis and its accuracy is demanded by the developer as well. Our work involved smart contracts with well-known vulnerabilities, but the definition of the properties is an indispensable and complex step. The tools without specification implement vulnerabilities detectors, which are stated in their paper. They have well-defined limitations, but strengths as well. These detect preconfigured vulnerabilities, so a warning can be run even if we expect that indeed the number of false negatives can be relevant. An example is Slither, which for every comparison of block time stamp give a warning. These warnings

Our results stress the impossibility of the tools without specification for the detection of specific vulnerabilities, but they could correctly detect the reentrancy cases. On the other hand, the ones with specifications had problems in the reentrancy detection. The limitations of some of these involve the external calls. SolcVerify was the only tool which could provide the possibility of reentrancy detection. Echidna, as Certora, developers teams specified the tools can detect the reentrancy, in the case an attack is provided, but we think this approach can be useful for checking a possible attack rather than detection of vulnerabilities.

Manticore could bridge this gap by adopting two different running modes, so the user, knowing the limitation of each way of analyses, can combine those for obtaining a valuable result.

7.3 Combination of Tools

8 Conclusion

General overview to the research questions, new approaches like Machine Learning, citing testing environment as Harhat, future of blockchain and adoption of solidity, different vulnerabilities regarding even the architecture itself. 1 or 2 pages

A Appendix

A.1 First Appendix Section

Figure A.1: A figure

...